

Mutual Exclusion (Mutex)

Protects critical sections that are longer than 1 line

Mutual Exclusion (Mutex)

- Enforcement
 - Only **1** thread in a critical section at a time
- Availability
 - If **no** thread in critical section, then **any** thread can enter
- Minimal Stay
 - Threads stay in critical section for minimal time
- Consistency
 - If resource must be protected anywhere, then it must be protected everywhere

Mutex – how to implement

- Hardware enforced
 - Disable interrupts
 - Guarantees atomic code because your code cannot be interrupted

Mutex – how to implement

- Hardware enforced
 - Disable interrupts
 - Guarantees atomic code because your code cannot be interrupted
- But...
 - Cannot have overlapping critical sections
 - Cannot switch to other, non-related, processes
 - Will not work on multi-core system unless you disable interrupts on all cores (big performance hit)
 - Kills performance on core

Mutex – how to implement

- Hardware enforced
 - Disable interrupts
 - Guarantees atomic code because your code cannot be interrupted
- But...
 - Cannot have overlapping critical sections
 - Cannot switch to other, non-related, processes
 - Will not work on multi-core system unless you disable interrupts on all cores (big performance hit)
 - Kills performance on core
- So... Cannot use disabled interrupts solution

Mutex – how to implement

```
1 //a special atomic function
2 int compare_and_swap(int *word, int notlockedval, int lockedval){
3     //save original
4     int oldval=*word;
5
6     if(oldval==notlockedval)    //if we can
7     |   *word=lockedval        //then do
8     return oldval;
9 }
10
11 const int notlockedval=0;
12 const int lockedval=1;
13
14 int word=notlockedval; //start unlocked
15
16 void withdraw(int amt){
17
18     //stay in below loop until compare_and_swap returns notlockedval
19     while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20     if(balance > amt){
21         cout<<"approved"<<endl;
22         balance-=amount;
23     }
24     word= notlockedval;
25 }
```

Compare and Swap – an atomic operation
(once started cannot be interrupted)

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |   *word=lockedval        //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap – an atomic operation
(once started cannot be interrupted)

1st thread => word=notlockedval => line 19 returns false, go to 20

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |   *word=lockedval        //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap – an atomic operation
(once started cannot be interrupted)

1st thread => word=notlockedval => line 19 returns false, go to 20

1st thread => Interrupted at line 21

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |   *word=lockedval         //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap – an atomic operation
(once started cannot be interrupted)

1st thread => word=notlockedval => line 19 returns false, go to 20

1st thread => Interrupted at line 21

2nd thread => line 19 returns true, so it busy waits (spins) burning CPU time

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |    *word=lockedval        //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap – an atomic operation
(once started cannot be interrupted)

1st thread => word=notlockedval => line 19 returns false, go to 20

1st thread => Interrupted at line 21

2nd thread => line 19 returns true, so it busy waits (spins) burning CPU time

1st thread swapped back and finishes, line 24 set word= word=notlockedval

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |    *word=lockedval        //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap – an atomic operation
(once started cannot be interrupted)

1st thread => word=notlockedval => line 19 returns false, go to 20

1st thread => Interrupted at line 21

2nd thread => line 19 returns true, so it busy waits (spins) burning CPU time

1st thread swapped back and finishes, line 24 set word= word=notlockedval

2nd thread swapped back, line 19 returns false, goes to 20 and finishes

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |   *word=lockedval        //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap – an atomic operation (once started cannot be interrupted)

1st thread => word=notlockedval => line 19 returns false, go to 20

1st thread => Interrupted at line 21

2nd thread => line 19 returns true, so it busy waits (spins) burning CPU time

1st thread swapped back and finishes, line 24 set word= word=notlockedval

2nd thread swapped back, line 19 returns false, goes to 20 and finishes

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |    *word=lockedval        //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap

Good:

Simple

Easily verified

Multiprocessor/multiprocess as long as can share memory

Mutex – how to implement

```
1  //a special atomic function
2  int compare_and_swap(int *word, int notlockedval, int lockedval){
3      //save original
4      int oldval=*word;
5
6      if(oldval==notlockedval)    //if we can
7      |    *word=lockedval        //then do
8      return oldval;
9  }
10
11  const int notlockedval=0;
12  const int lockedval=1;
13
14  int word=notlockedval; //start unlocked
15
16  void withdraw(int amt){
17
18      //stay in below loop until compare_and_swap returns notlockedval
19      while (compare_and_swap(&word,notlockedval, lockedval ) == lockedval){}
20      if(balance > amt){
21          cout<<"approved"<<endl;
22          balance-=amount;
23      }
24      word= notlockedval;
25  }
```

Compare and Swap

Good:

Simple

Easily verified

Multiprocessor/multiprocess as long as can share memory

Bad

Busy wait (line 19) must keep checking until available, CPU usage spikes)

Starvation and Deadlock both possible

Compare_and_swap has to be atomic, this C++ code is not

Mutex – Using implementation in C++ 11

- Mutexes are thread based in C++ 11, not process based!

```
#include <mutex>
```

```
std::mutex g_mutex;          //generally a global value
```

```
:
```

```
g_mutex.lock();              //if available then proceed, otherwise thread blocks
```

```
g_mutex.unlock(); //unlocks the mutex, other waiting threads can acquire
```

General rules

Unlock a mutex when you are done (else waiting threads will wait forever)

Do not lock() a mutex twice from same thread without intermediate unlock(). Otherwise thread will block waiting to acquire a mutex that it has.

Mutex – Solve withdrawal problem

```
std::mutex mymutex;  
void withdraw(int amt){  
    mymutex.lock();  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    mymutex.unlock();  
}
```


Mutex – Solve withdrawal problem

```
std::mutex mymutex;  
void withdraw(int amt){  
    mymutex.lock();  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    mymutex.unlock();  
}
```

Mutex – Solve withdrawal problem

```
std::mutex mymutex;  
void withdraw(int amt){  
    mymutex.lock();  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    mymutex.unlock();  
}
```

But wait! What if you throw an exception here?



Mutex – Solve withdrawal problem

```
std::mutex mymutex;  
void withdraw(int amt){  
    mymutex.lock();  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    mymutex.unlock();  
}
```

But wait! What if you throw an exception here?

- You will never unlock the mutex.
- All threads waiting to enter the critical section will be blocked forever
- Process will never join() those threads
- Process will be blocked forever
- Have to kill and restart process

Mutex – A better idea

- Use a self unlocking mutex- As soon as the mutex goes out of scope it unlocks.

```
std::mutex mymutex;  
void withdraw(int amt){  
    lock_guard<std::mutex> lock(mymutex); //locks mymutex here  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    //unlocks mymutex here when the lock_guard goes out of scope  
}  
  
lock_guard<std::mutex> lock(mymutex);
```

Mutex – A better idea

- Use a self unlocking mutex- As soon as the mutex goes out of scope it unlocks.

```
std::mutex mymutex;  
void withdraw(int amt){  
    lock_guard<std::mutex> lock(mymutex); //locks mymutex here  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    //unlocks mymutex here when the lock_guard goes out of scope  
}  
  
lock_guard<std::mutex> lock(mymutex);
```

Mutex – A better idea

- Use a self unlocking mutex- As soon as the mutex goes out of scope it unlocks.

```
std::mutex mymutex;  
void withdraw(int amt){  
  
    lock_guard<std::mutex> lock(mymutex); //locks mymutex here  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    //unlocks mymutex here when the lock_guard goes out of scope  
}
```

```
lock_guard<std::mutex> lock(mymutex);
```

But wait! What if you throw an exception here?



Mutex – A better idea

- Use a self unlocking mutex- As soon as the mutex goes out of scope it unlocks.

```
std::mutex mymutex;  
void withdraw(int amt){  
  
    lock_guard<std::mutex> lock(mymutex); //locks mymutex here  
    if(balance > amt){  
        cout<<"approved"<<endl;  
        balance-=amount;  
    }  
    //unlocks mymutex here when the lock_guard goes out of scope  
}
```

```
lock_guard<std::mutex> lock(mymutex);
```

But wait! What if you throw an exception here?

- No worries, the lock_guard will unlock as soon as it goes out of scope